

Media_Mix_Model

June 16, 2025

0.1 Importing Required Libraries

These libraries are necessary for data manipulation, visualization, and statistical modeling.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import statsmodels.api as sm
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.stats.diagnostic import acorr_ljungbox, het_breuschpagan
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

0.2 Exploratory Data Analysis (EDA)

This section explores the dataset's structure, content, and variable distributions.

0.3 Data Preparation Summary

This step involves loading the dataset, checking for missing or duplicate values, and understanding the basic structure.

0.3.1 1. Loading Data and data description

```
[2]: data = pd.read_csv("../data/data.csv")
```

```
[3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6240 entries, 0 to 6239
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   geo                  6240 non-null   object
1   date                 6240 non-null   object
2   tv_impression        6240 non-null   int64
3   radio_impression     6240 non-null   int64
4   print_impression     6240 non-null   int64
5   search_impression    6240 non-null   int64
6   social_impression    6240 non-null   int64
```

```

7 competitor_sales_control 6240 non-null float64
8 sentiment_score_control 6240 non-null float64
9 tv_spend 6240 non-null float64
10 radio_spend 6240 non-null float64
11 print_spend 6240 non-null float64
12 search_spend 6240 non-null float64
13 social_spend 6240 non-null float64
14 population 6240 non-null float64
15 sales 6240 non-null float64

```

dtypes: float64(9), int64(5), object(2)

memory usage: 780.1+ KB

[4]: data.describe()

```

[4]:      tv_impression  radio_impression  print_impression  search_impression \
count  6.240000e+03    6.240000e+03    6.240000e+03    6.240000e+03
mean   8.851635e+05    5.206055e+05    2.605050e+05    1.806810e+06
std    8.193451e+05    6.206231e+05    5.551842e+05    1.304818e+06
min    0.000000e+00    0.000000e+00    0.000000e+00    0.000000e+00
25%    2.536458e+05    0.000000e+00    0.000000e+00    8.083565e+05
50%    6.796705e+05    3.106525e+05    0.000000e+00    1.505379e+06
75%    1.324174e+06    8.075525e+05    2.547340e+05    2.566118e+06
max    5.192032e+06    4.397540e+06    5.610156e+06    7.635147e+06

      social_impression  competitor_sales_control  sentiment_score_control \
count  6.240000e+03    6240.000000    6240.000000
mean   9.816780e+05    -0.044412    -0.033675
std    9.141768e+05    1.210164    1.175572
min    0.000000e+00    -4.825634    -4.230392
25%    2.675945e+05    -0.872736    -0.808085
50%    7.462475e+05    -0.050346    -0.036433
75%    1.476708e+06    0.743633    0.739391
max    6.975542e+06    3.924682    4.320259

      tv_spend  radio_spend  print_spend  search_spend  social_spend \
count  6240.000000  6240.000000  6240.000000  6240.000000  6240.000000
mean   6490.659147  5019.278558  1935.793164  14080.185617  7649.147993
std    6008.031135  5983.571424  4125.531614  10168.239254  7123.184932
min    0.000000    0.000000    0.000000    0.000000    0.000000
25%    1859.914025  0.000000    0.000000    6299.393250  2085.072625
50%    4983.835500  2995.072650    0.000000  11731.178500  5814.694600
75%    9709.803000  7785.800675  1892.909150  19997.351500  11506.376500
max    38071.730000  42397.700000  41688.645000  59499.484000  54352.810000

      population  sales
count  6240.000000  6240.000000
mean   542418.832750  211330.038098

```

```

std    242839.719887  121605.908396
min    136670.940000   9742.282233
25%    335176.345000  114323.174425
50%    560478.825000  193995.582650
75%    736033.810000  287918.213625
max    994048.940000  739823.338900

```

```
[5]: data.head()
```

```

[5]:      geo      date  tv_impression  radio_impression  print_impression  \
0  Geo0  1/25/21      280668             0                0
1  Geo0  2/1/21      366206            182108            19825
2  Geo0  2/8/21      197565            230170                0
3  Geo0  2/15/21     140990            66643                0
4  Geo0  2/22/21     399116            164991                0

      search_impression  social_impression  competitor_sales_control  \
0              470611             108010             -1.338765
1              527702             252506              0.893645
2              393618             184061             -0.284549
3              326034             201729             -1.034740
4              381982             153973             -0.319276

      sentiment_score_control  tv_spend  radio_spend  print_spend  search_spend  \
0              0.115581  2058.0608      0.00000      0.00000      3667.3965
1              0.944224  2685.2874  1755.74540  147.31808      4112.2974
2             -1.290579  1448.6895  2219.12230      0.00000      3067.4023
3             -1.084514  1033.8406   642.52057      0.00000      2540.7310
4             -0.017503  2926.6072  1590.71640      0.00000      2976.7249

      social_spend  population      sales
0      841.6044   136670.94  39198.55690
1     1967.5044   136670.94  41497.96063
2     1434.1870   136670.94  41579.08885
3     1571.8545   136670.94  56492.86151
4     1199.7440   136670.94  71039.82718

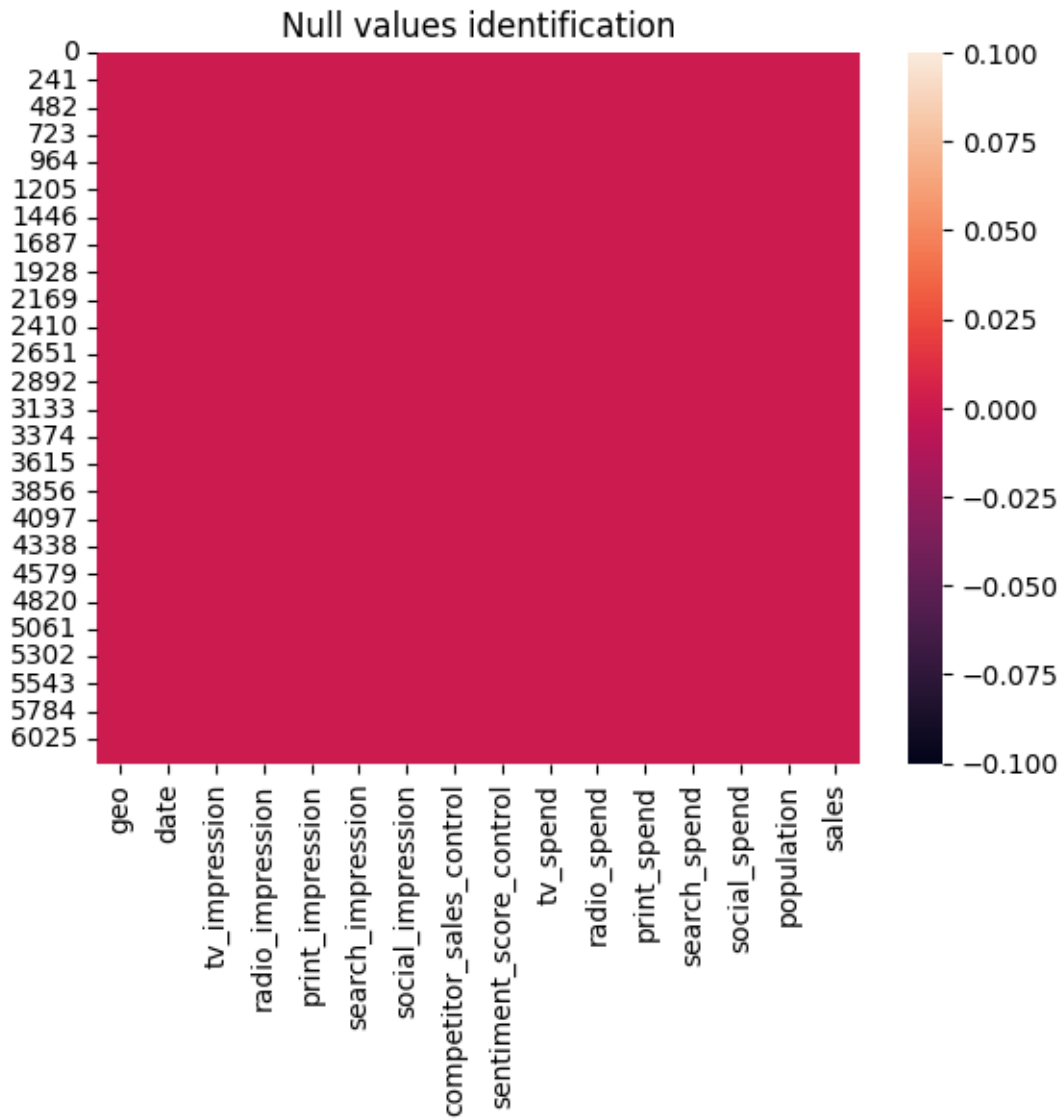
```

0.3.2 2. Null values identification

```

[6]: sns.heatmap(data.isnull())
plt.title("Null values identification")
plt.show()

```



```
[7]: print(f"Number of null values: {data.isnull().sum().sum()}")
```

Number of null values: 0

0.3.3 3. Duplicates detection

```
[8]: duplicates = data[data.duplicated()]
print(f"Number of duplicates records: {len(duplicates)}")
```


Number of duplicates records: 0

0.3.4 1. Pearson Correlation Heatmap

Visualizes linear relationships between numerical variables using Pearson's method.

```
[9]: correlations = data.select_dtypes(include=['number']).corr()

sns.heatmap(correlations)
plt.title("Features correlation matrix")
plt.show()
```



0.3.5 2. Correlation with Target (Sales)

Highlights which variables are most strongly correlated (positively and negatively) with sales.

```
[10]: # Compute correlations with 'sales'
numeric_cols = data.select_dtypes(include=['number'])
correlations = numeric_cols.corr()['sales'].drop('sales')

# Top 5 positive and negative correlations
top_positive = correlations.sort_values(ascending=False).head(5).reset_index()
top_negative = correlations.sort_values(ascending=True).head(5).reset_index()
```

```

# Rename columns
top_positive.columns = ['variable', 'correlation']
top_negative.columns = ['variable', 'correlation']

# Create figure
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(14, 5))

# Plot positive correlations
sns.barplot(
    data=top_positive,
    x='correlation', y='variable',
    ax=axes[0], palette='Greens_r'
)
axes[0].set_title("Top 5 Positive Correlations")
axes[0].set_xlabel("Correlation")
axes[0].set_ylabel("Variable")

# Plot negative correlations
sns.barplot(
    data=top_negative,
    x='correlation', y='variable',
    ax=axes[1], palette='Reds'
)
axes[1].set_title("Top 5 Negative Correlations")
axes[1].set_xlabel("Correlation")
axes[1].set_ylabel("Variable")

plt.tight_layout()
plt.show()

```

C:\Users\Rafa Rincon\AppData\Local\Temp\ipykernel_18392\1830579637.py:17:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

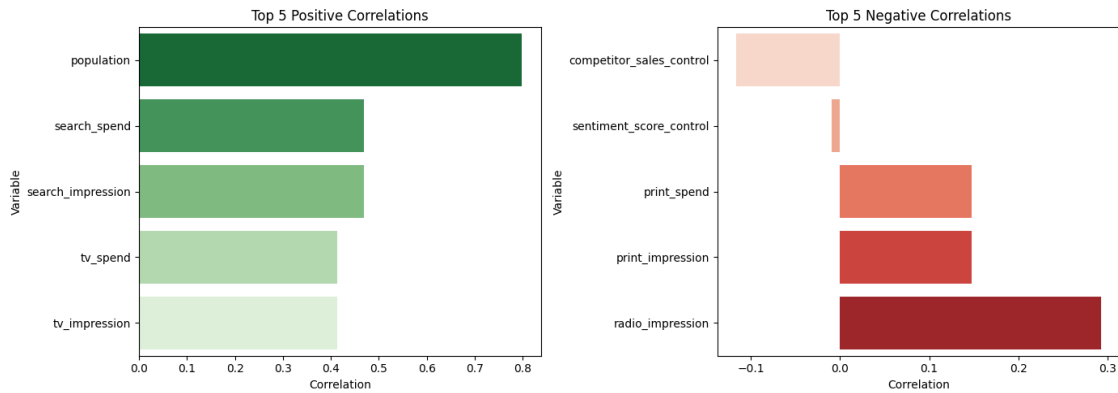
```
sns.barplot(
```

C:\Users\Rafa Rincon\AppData\Local\Temp\ipykernel_18392\1830579637.py:27:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



0.4 Feature Engineering

New features are derived from existing data to enhance model performance and capture relevant signals.

0.4.1 1. Geographic Feature Encoding

Converts the 'geo' categorical variable into binary (dummy) variables for use in regression models.

```
[11]: # Convert the 'geo' categorical variable into dummy/indicator variables (drop
      ↳ the first to avoid multicollinearity)
      geo_dummies = pd.get_dummies(data['geo'], drop_first=True)

      # Ensure all dummy variables are of integer type
      geo_dummies = geo_dummies.astype(int)

      # Append the dummy variables to the original dataset
      data = pd.concat([data, geo_dummies], axis=1)
```

0.4.2 2. Temporal Feature Transformation

Extracts seasonality and calendar-related components from the date column.

```
[12]: data['date'] = pd.to_datetime(data['date'])
      data['year'] = data['date'].dt.year
      data['month'] = data['date'].dt.month
      data['weekofyear'] = data['date'].dt.isocalendar().week.astype(int)

      #data['is_weekend'] = data['date'].dt.dayofweek >= 5
      #data['is_weekend'] = data['is_weekend'].astype(int)
      data['quarter'] = data['date'].dt.quarter
      data['month_sin'] = np.sin(2 * np.pi * data['month'] / 12)
      data['month_cos'] = np.cos(2 * np.pi * data['month'] / 12)
```

```
C:\Users\Rafa Rincon\AppData\Local\Temp\ipykernel_18392\1900084449.py:1:
UserWarning: Could not infer format, so each element will be parsed
individually, falling back to `dateutil`. To ensure parsing is consistent and
as-expected, please specify a format.
    data['date'] = pd.to_datetime(data['date'])
```

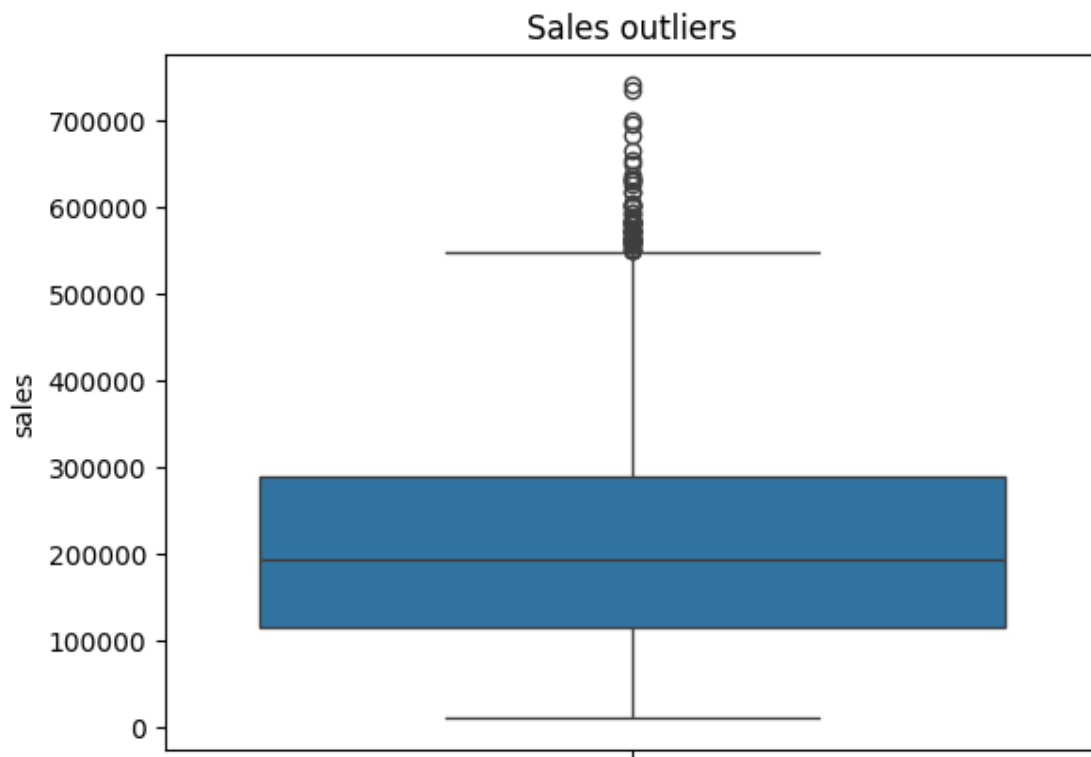
0.5 Seasonality and Time-Series Structure

Analyzes temporal sales patterns, decomposing them into trend, seasonal, and irregular components.

0.5.1 1. Outlier Detection in Sales

Boxplots help identify extreme values or irregular fluctuations in the sales data.

```
[13]: sns.boxplot(data["sales"])
plt.title("Sales outliers")
plt.show()
```



0.5.2 2. Time-Series Decomposition

Uses additive decomposition to separate sales into seasonal, trend, and residual components.


```
[14]: # Filter data by a specific geography and aggregate sales by date
geo_df = data[data['geo'] == 'Geo0'].groupby('date')['sales'].sum()
geo_df = geo_df.sort_index()

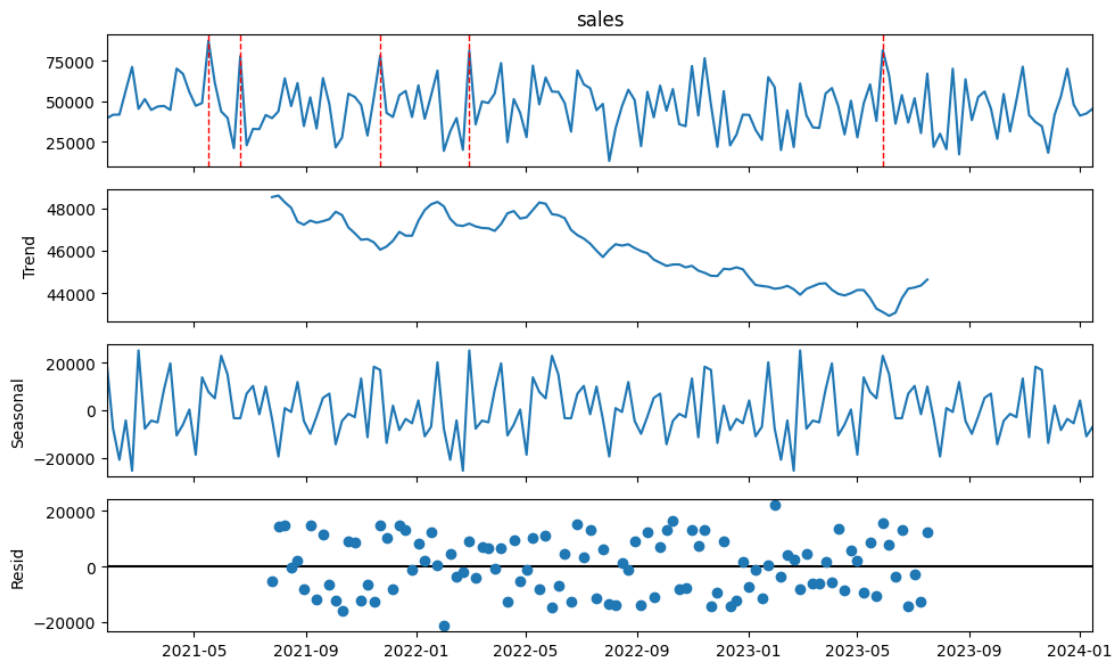
# Perform seasonal decomposition (weekly frequency, assuming 52 periods per
# year)
result = seasonal_decompose(geo_df, model='additive', period=52)

# Identify top 5 dates with the highest sales (sales peaks)
top_peaks = geo_df.sort_values(ascending=False).head(5).index

# Plot decomposition components and mark sales peaks with red lines
fig = result.plot()
fig.set_size_inches(10, 6)

# Add vertical red dashed lines on the observed component for peak dates
observed_ax = fig.axes[0]
for date in top_peaks:
    observed_ax.axvline(x=date, color='red', linestyle='--', linewidth=1)

plt.tight_layout()
plt.show()
```



0.6 Media Mix Modeling: Channel Effectiveness

This model evaluates how different marketing channels contribute to overall sales.

0.6.1 1. Adstock and Saturation Effects

Applies an adstock transformation to simulate delayed media effects, followed by log transformation for diminishing returns.

```
[15]: # Define a function to apply the adstock transformation
# This simulates the carryover effect of advertising over time using a decay
      ↪ factor
def apply_adstock(series, decay=0.5):
    result = []
    for i, val in enumerate(series):
        if i == 0:
            result.append(val) # First value remains unchanged
        else:
            # Current value plus a portion of the previous value based on the
            ↪ decay factor
            result.append(val + decay * result[i - 1])
    return pd.Series(result, index=series.index)

# List of media spend channels to apply adstock to
spend_channels = ['tv_spend', 'radio_spend', 'print_spend', 'search_spend',
                  ↪ 'social_spend']

# Set decay rate for adstock effect
decay = 0.5

# Apply adstock and log transformation (to model diminishing returns) for each
      ↪ channel
for channel in spend_channels:
    adstocked = apply_adstock(data[channel], decay)
    data[f'{channel}_transformed'] = np.log1p(adstocked) # Log transformation
    ↪ for saturation effect
```

0.6.2 2. Regression Coefficient Interpretation

Estimates the marginal effect of each feature on sales using Ordinary Least Squares (OLS) regression.

```
[16]: features_adstock = [
    "tv_spend_transformed", "radio_spend_transformed",
    ↪ "print_spend_transformed",
    "search_spend_transformed", "social_spend_transformed",
    'competitor_sales_control', 'sentiment_score_control', 'population',
    'month_sin', 'month_cos'
]

# Add geographic dummy variables (geo_*) to the list of features
geo_features = geo_dummies.columns.tolist()
```

```
features_with_geo = features_adstock + geo_features
```

```
[17]: # Add a constant (intercept) term to the feature set
X_with_geo = sm.add_constant(data[features_with_geo])

# Define the target variable (sales)
y_with_geo = data['sales']

# Fit an Ordinary Least Squares (OLS) regression model using the selected
↪ features
model_with_geo = sm.OLS(y_with_geo, X_with_geo).fit()

# Display the summary of regression results
print(model_with_geo.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          sales    R-squared:                0.696
Model:                  OLS      Adj. R-squared:           0.694
Method:                 Least Squares    F-statistic:           295.2
Date:                  Mon, 16 Jun 2025    Prob (F-statistic):       0.00
Time:                  11:06:03    Log-Likelihood:          -78201.
No. Observations:      6240    AIC:                    1.565e+05
Df Residuals:          6191    BIC:                    1.568e+05
Df Model:              48
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025
0.975]					
const	-1.526e+05	3.08e+04	-4.960	0.000	-2.13e+05
-9.23e+04					
tv_spend_transformed	4379.7523	1877.944	2.332	0.020	698.330
8061.175					
radio_spend_transformed	2549.3814	1227.115	2.078	0.038	143.810
4954.953					
print_spend_transformed	1553.4649	547.205	2.839	0.005	480.752
2626.177					
search_spend_transformed	7195.1300	3424.370	2.101	0.036	482.176
1.39e+04					
social_spend_transformed	2964.4336	2012.665	1.473	0.141	-981.089
6909.956					
competitor_sales_control	-1.517e+04	882.308	-17.194	0.000	-1.69e+04
-1.34e+04					
sentiment_score_control	-4354.0644	879.537	-4.950	0.000	-6078.263
-2629.866					

population	0.3489	0.015	23.651	0.000	0.320
0.378					
month_sin	-4076.6565	1380.674	-2.953	0.003	-6783.257
-1370.056					
month_cos	1775.7455	1330.578	1.335	0.182	-832.649
4384.140					
Geo1	-7430.5528	7241.948	-1.026	0.305	-2.16e+04
6766.180					
Geo10	1.509e+04	6295.585	2.396	0.017	2745.811
2.74e+04					
Geo11	-3.007e+04	5323.410	-5.649	0.000	-4.05e+04
-1.96e+04					
Geo12	-9183.6249	5509.839	-1.667	0.096	-2e+04
1617.573					
Geo13	-5776.0835	5992.854	-0.964	0.335	-1.75e+04
5971.992					
Geo14	-2.647e+04	5345.552	-4.951	0.000	-3.69e+04
-1.6e+04					
Geo15	1.283e+04	5312.333	2.415	0.016	2412.962
2.32e+04					
Geo16	-6194.2840	6324.056	-0.979	0.327	-1.86e+04
6203.062					
Geo17	-2.007e+04	5364.338	-3.741	0.000	-3.06e+04
-9554.135					
Geo18	-2.96e+04	5415.389	-5.466	0.000	-4.02e+04
-1.9e+04					
Geo19	-2.002e+04	6109.157	-3.278	0.001	-3.2e+04
-8047.415					
Geo2	-3.12e+04	5690.988	-5.483	0.000	-4.24e+04
-2e+04					
Geo20	-1.307e+04	5528.165	-2.364	0.018	-2.39e+04
-2231.386					
Geo21	-3.822e+04	5511.291	-6.935	0.000	-4.9e+04
-2.74e+04					
Geo22	7690.0288	6124.062	1.256	0.209	-4315.259
1.97e+04					
Geo23	-1.494e+04	6227.989	-2.399	0.016	-2.71e+04
-2729.040					
Geo24	1894.3192	7599.960	0.249	0.803	-1.3e+04
1.68e+04					
Geo25	9408.5695	5365.957	1.753	0.080	-1110.569
1.99e+04					
Geo26	-7649.6668	6865.332	-1.114	0.265	-2.11e+04
5808.767					
Geo27	-7482.6548	5529.120	-1.353	0.176	-1.83e+04
3356.340					
Geo28	7049.2178	5308.324	1.328	0.184	-3356.941
1.75e+04					

Geo29	1.188e+04	6042.941	1.966	0.049	33.547
2.37e+04					
Geo3	-4606.0132	7071.983	-0.651	0.515	-1.85e+04
9257.528					
Geo30	7080.0069	6838.591	1.035	0.301	-6326.006
2.05e+04					
Geo31	3710.7586	5656.443	0.656	0.512	-7377.834
1.48e+04					
Geo32	-2.056e+04	5310.595	-3.871	0.000	-3.1e+04
-1.01e+04					
Geo33	-7233.0453	5458.467	-1.325	0.185	-1.79e+04
3467.446					
Geo34	-5059.8442	5405.985	-0.936	0.349	-1.57e+04
5537.763					
Geo35	-2.481e+04	5310.746	-4.671	0.000	-3.52e+04
-1.44e+04					
Geo36	4.657e+04	6166.777	7.551	0.000	3.45e+04
5.87e+04					
Geo37	3020.6015	6563.431	0.460	0.645	-9846.003
1.59e+04					
Geo38	3477.4263	7016.740	0.496	0.620	-1.03e+04
1.72e+04					
Geo39	1.026e+05	5423.319	18.909	0.000	9.19e+04
1.13e+05					
Geo4	8606.5773	6474.249	1.329	0.184	-4085.200
2.13e+04					
Geo5	-1.203e+04	6473.701	-1.858	0.063	-2.47e+04
665.616					
Geo6	-1475.8298	7128.859	-0.207	0.836	-1.55e+04
1.25e+04					
Geo7	4.445e+04	5315.113	8.362	0.000	3.4e+04
5.49e+04					
Geo8	3900.6397	6645.157	0.587	0.557	-9126.175
1.69e+04					
Geo9	4.031e+04	5695.099	7.079	0.000	2.91e+04
5.15e+04					

```
=====
Omnibus:                221.064    Durbin-Watson:                1.988
Prob(Omnibus):           0.000    Jarque-Bera (JB):          646.214
Skew:                    -0.003    Prob(JB):                  4.75e-141
Kurtosis:                 4.577    Cond. No.                   3.99e+21
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.38e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

0.6.3 3. ROI and Marginal Contribution Analysis

Calculates Return on Investment (ROI) and average contribution to sales by media channel.

```
[18]: # Define the list of original media spend channels
spend_channels = ['tv_spend', 'radio_spend', 'print_spend', 'search_spend',
                 ↪ 'social_spend']

# Retrieve model coefficients (from the geo-inclusive model or use
↪ model_adstock if geo is not included)
coeffs = model_with_geo.params

# Calculate contribution and ROI for each media channel
contrib_values = []
roi_values = []

for ch in spend_channels:
    ch_trans = f'{ch}_transformed' # Transformed version of the spend variable
    coef = coeffs[ch_trans] # Coefficient from the regression model
    avg_trans = data[ch_trans].mean() # Average transformed spend
    avg_spend = data[ch].mean() # Average actual spend

    # Contribution is the coefficient times the average transformed value
    contribution = coef * avg_trans

    # ROI is contribution divided by the average actual spend
    roi = contribution / avg_spend if avg_spend != 0 else None

    # Store results
    contrib_values.append(contribution)
    roi_values.append(roi)

# Create a DataFrame to summarize contributions and ROIs by channel
df_contrib_roi_geo = pd.DataFrame({
    'channel': spend_channels,
    'contribution': contrib_values,
    'roi': roi_values
})
```

```
[19]: df_contrib_roi_geo
```

```
[19]:
```

	channel	contribution	roi
0	tv_spend	40426.675355	6.228439
1	radio_spend	22575.263459	4.497711
2	print_spend	10995.905744	5.680310
3	search_spend	72374.227808	5.140147
4	social_spend	27812.133260	3.635978

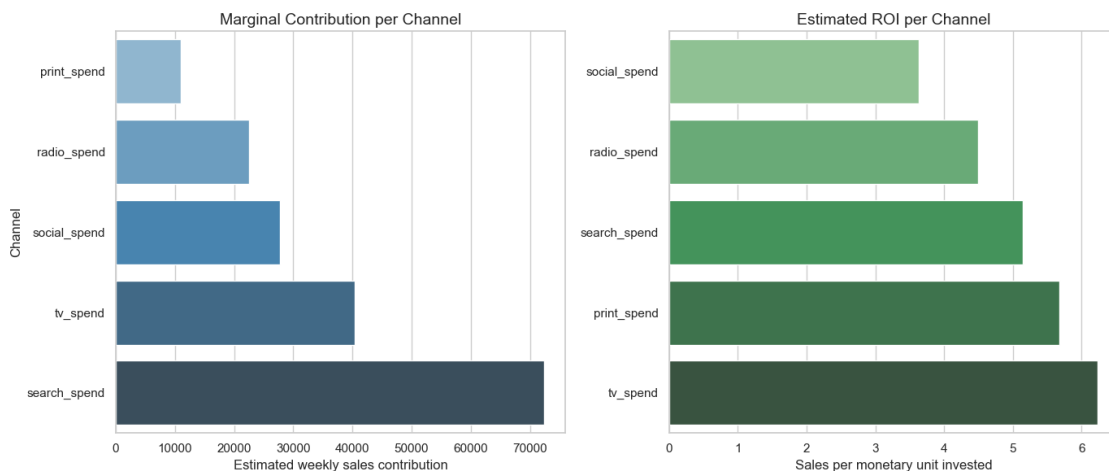
```
[20]: # Set visual style
sns.set(style="whitegrid")

# Create figure with two horizontal bar plots
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Contribution plot
sns.barpplot(
    x="contribution", y="channel", hue="channel", legend=False,
    data=df_contrib_roi_geo.sort_values("contribution", ascending=True),
    palette="Blues_d", ax=axes[0]
)
axes[0].set_title("Marginal Contribution per Channel", fontsize=14)
axes[0].set_xlabel("Estimated weekly sales contribution")
axes[0].set_ylabel("Channel")

# ROI plot
sns.barpplot(
    x="roi", y="channel", hue="channel", legend=False,
    data=df_contrib_roi_geo.sort_values("roi", ascending=True),
    palette="Greens_d", ax=axes[1]
)
axes[1].set_title("Estimated ROI per Channel", fontsize=14)
axes[1].set_xlabel("Sales per monetary unit invested")
axes[1].set_ylabel("") # Hide redundant label

# Adjust layout for clarity
plt.tight_layout()
plt.show()
```



0.6.4 4. Regional and Temporal ROI Analysis

Examines how ROI varies across geographies and time periods to uncover channel performance patterns.

```
[21]: # Define the list of media spend channels
spend_channels = ['tv_spend', 'radio_spend', 'print_spend', 'search_spend',
                  ↪ 'social_spend']

# Extract global model coefficients
coeffs = model_with_geo.params

# Initialize list to store ROI and contribution results by region and time
geo_time_results = []

# Group the data by geography, year, month, and quarter
for (geo, year, month, quarter), group in data.groupby(['geo', 'year', 'month',
                  ↪ "quarter"]):
    for ch in spend_channels:
        ch_trans = f'{ch}_transformed'
        coef = coeffs.get(ch_trans, 0) # Get coefficient; use 0 if not found

        avg_trans = group[ch_trans].mean() # Mean transformed media spend for
        ↪ the group
        avg_spend = group[ch].mean() # Mean actual media spend for the
        ↪ group

        contribution = coef * avg_trans # Compute estimated contribution to
        ↪ sales
        roi = contribution / avg_spend if avg_spend != 0 else None # Compute
        ↪ ROI

    # Store the result for this geo-time-channel combination
    geo_time_results.append({
        'geo': geo,
        'year': year,
        'month': month,
        "quarter": quarter,
        'channel': ch,
        'contribution': contribution,
        'roi': roi
    })

# Create a final DataFrame with the results
df_geo_month_roi = pd.DataFrame(geo_time_results)

# Export results to CSV (optional)
```



```
df_geo_month_roi.to_csv("../data/ROI_Contribution_geo_year.csv", index=False)

# Display the top results sorted by ROI in descending order
print(df_geo_month_roi.sort_values(by="roi", ascending=False).head())
```

	geo	year	month	quarter	channel	contribution	roi
7332	Geo9	2022	12	4	print_spend	10544.370508	2752.638300
752	Geo12	2021	3	1	print_spend	7042.527624	1362.467476
57	Geo0	2021	12	4	print_spend	7497.059152	1122.246918
4952	Geo32	2023	5	2	print_spend	3273.768886	896.540688
4452	Geo30	2021	3	1	print_spend	8517.333353	740.344313

0.7 Model Evaluation

This section includes statistical validation of the regression model using residual diagnostics.

0.7.1 1. Residual Diagnostics

Uses the Ljung-Box and Breusch-Pagan tests to check for autocorrelation and heteroskedasticity in residuals.

```
[22]: # Extract model residuals and fitted (predicted) values
residuals = model_with_geo.resid
fitted = model_with_geo.fittedvalues

# Plot residuals vs. fitted values to visually assess homoscedasticity
↳ (constant variance)
plt.figure(figsize=(8, 5))
plt.scatter(fitted, residuals, alpha=0.5)
plt.axhline(0, color='red', linestyle='--') # Reference line at zero
plt.title("Residuals vs. Fitted Values")
plt.xlabel("Fitted Values")
plt.ylabel("Residuals")
plt.tight_layout()
plt.show()

# Plot Autocorrelation Function (ACF) and Partial Autocorrelation Function
↳ (PACF) of residuals
fig, ax = plt.subplots(1, 2, figsize=(14, 5))
plot_acf(residuals, lags=30, ax=ax[0]) # Autocorrelation plot
plot_pacf(residuals, lags=30, ax=ax[1]) # Partial autocorrelation plot
ax[0].set_title("ACF of Residuals")
ax[1].set_title("PACF of Residuals")
plt.tight_layout()
plt.show()

# Perform Ljung-Box test to check for autocorrelation in residuals
ljung_box = acorr_ljungbox(residuals, lags=[12], return_df=True)
```

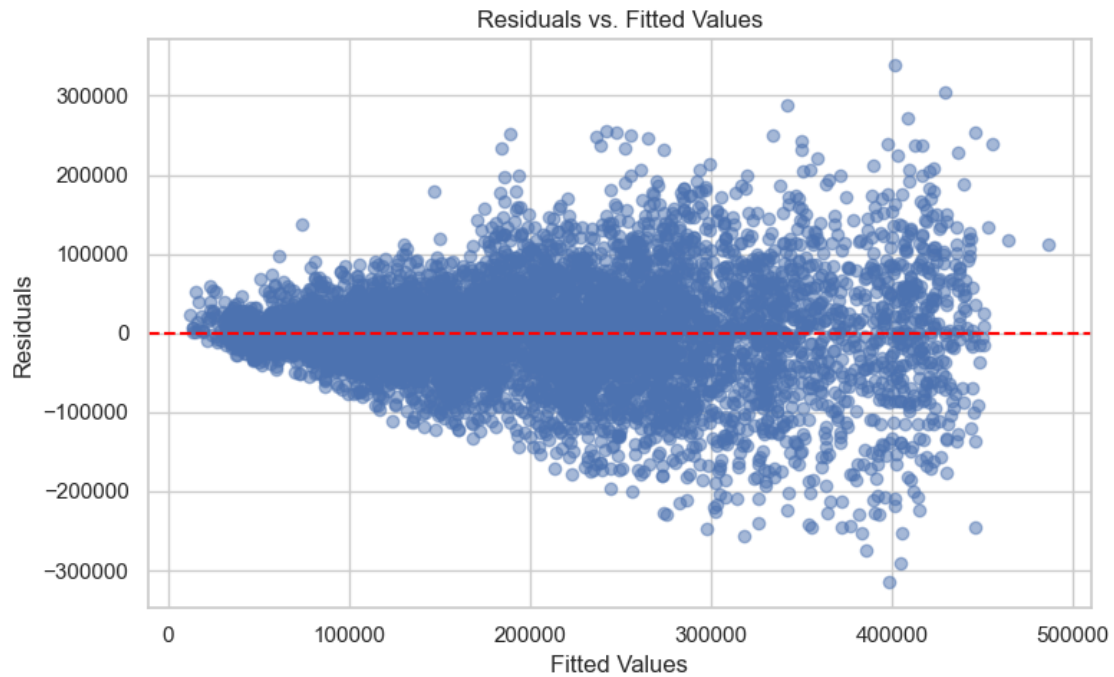
```

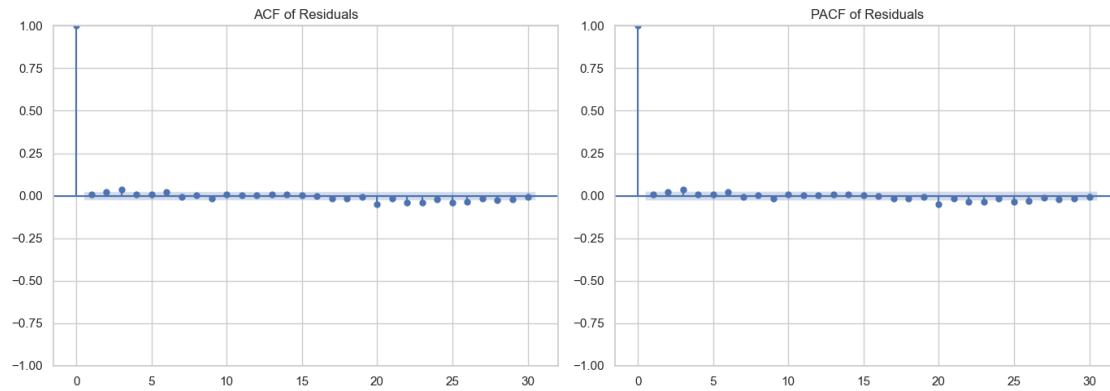
print("Ljung-Box test (lag=12):")
print(ljung_box)

# Perform Breusch-Pagan test to detect heteroskedasticity (non-constant
  ↳ variance of residuals)
exog = model_with_geo.model.exog
bp_test = het_breuschpagan(residuals, exog)

# Output Breusch-Pagan test results
print("\nBreusch-Pagan test:")
print(f"Test statistic: {bp_test[0]:.2f}")
print(f"p-value: {bp_test[1]:.4f}")
print(f"Interpretation: {'No heteroskedasticity detected' if bp_test[1] > 0.05
  ↳ else 'Heteroskedasticity present'}")

```





Ljung-Box test (lag=12):

	lb_stat	lb_pvalue
12	17.907369	0.118532

Breusch-Pagan test:
 Test statistic: 1027.96
 p-value: 0.0000
 Interpretation: Heteroskedasticity present